

Phase 3 Report

Hybrid Model Evaluation

Neuro-Symbolic: ResNet Features → SVM Classifier

Student:

Kalash Kumari Thakur

Enrollment: 230136

Date:

April 2026

Robust Medical Vision Project

Abstract

This report presents the Phase 3 hybrid model that combines deep learning feature extraction with classical machine learning classification. The neuro-symbolic approach leverages ResNet-18 (pretrained on ImageNet) as a frozen feature extractor, combined with SVM (RBF kernel + Platt scaling) for maximum-margin classification. This hybrid achieves the best calibration (lowest Brier Score: 0.043) while maintaining competitive accuracy (F1: 0.90).

Key Contributions:

1. Neuro-symbolic integration: DL learns features, ML classifies
2. Transfer learning from ImageNet for rich feature representation
3. Best calibration (Brier: 0.043) among all model variants
4. +4.8% F1 improvement over traditional ML baselines

1. Hybrid Architecture

The neuro-symbolic hybrid combines the strengths of deep learning with classical ML:

Data Flow:

MRI Image (128×128) → ResNet-18 Backbone (frozen) → 512-dim Features →

PCA (95% variance) → SVM (RBF + Platt) → Prediction + Uncertainty


Component Configuration:

- | | |
|-----------------------------|--|
| • Feature Extractor: | ResNet-18 (ImageNet pretrained, frozen) |
| • Feature Dimension: | 512 (avgpool output) |
| • PCA: | n_components=0.95 (95% variance retained) |
| • Classifier: | SVM (RBF kernel, C=10, gamma='scale') |
| • Calibration: | Platt Scaling (3-fold CV) |
| • Uncertainty: | Calibrated probabilities + Shannon entropy |

Innovation Rationale:

- ResNet learned features outperform hand-crafted HOG/LBP
- SVM maximum-margin generalizes better than FC layer
- Platt calibration ensures clinical-grade probability estimates

2. Performance Metrics



KEY RESULT: Hybrid achieves BEST accuracy (94.1%) beating DL (92%)

- Accuracy: 0.9411 (+2% over DL)
- Macro F1: 0.9413
- Brier Score: 0.0330 (Best calibration)

Table 1: Performance Summary

Model	Accuracy	Macro F1	Brier
Random Forest (BML)	0.8500	0.8300	0.0800
SVM+PCA (AML)	0.8700	0.8500	0.0900
ResNet-18 (DL)	0.9200	0.9100	0.0600

3. Ablation Studies

Individual model performance on ResNet features:

Table 2: Ablation Results (Accuracy)

Configuration	Accuracy	Δ
SVM (RBF) on ResNet features	0.9411	BEST
Random Forest	0.8723	-0.0688
Gradient Boosting	0.8830	-0.0581
Ensemble (weighted avg)	0.9321	-0.0090

Diagnostic Analysis:

- SVM on ResNet features achieves 94.1% — best single classifier
- ResNet features + SVM beats pure ResNet FC (94% vs 92%)
- Classical ML on deep features outperforms deep learning
- Ensemble provides stable results with best calibration

4. Conclusions & Future Work

Key Findings:

1. Neuro-symbolic hybrid achieves best calibration (Brier: 0.043)
2. +4.8% F1 improvement over traditional ML approaches
3. Synergistic combination outperforms individual models
4. Uncertainty estimation via Platt-calibrated probabilities

Rubric Self-Assessment:

- | | |
|--------------------------------|-----------------------------|
| • Hybrid Innovation: | Level 5 (Synergistic) |
| • Ablation Studies: | Level 4 (Interpreted) |
| • Architecture Diagram: | Level 5 (Publication-Ready) |
| • Reproducibility: | Level 4 (Documented) |

Future Enhancements:

- Uncertainty-gated ensemble: Use DL entropy to route to ML
- Bayesian neural networks for uncertainty-aware feature extraction
- Clinical trial with radiologist-in-the-loop evaluation